

Elf Formats in Linux

ELF: Executable and Linking Format

- ELF is a portable object file format defining the composition and organization of the object file. Kernel and binary loader looks at this format to know how to load this file and find various pieces of information like code, initialized data, dependencies on shared libraries etc.

Elf File Types

- Elf defines the format of executable binary files. There are four different types:
 - Relocatable
 - Created by compilers or assemblers. Need to be processed by the linker before running.
 - holds code and data suitable to link with other object files
 - Executable
 - Have all relocation done and all symbol resolved except perhaps shared library symbols that must be resolved at run time.
 - suitable for execution
 - Shared object
 - Shared library containing both symbol information for the linker and directly runnable code for run time.
 - holds code and data suitable to link with other relocatable object or shared objects
 - Core file
 - A core dump file.

ELF Structure

- Elf files have a dual nature:
 - Compilers, assemblers, and linkers treat the file as a set of logical sections described by a section header table.
 - The system loader treats the file as a set of segments described by a program header table.

ELF Structure

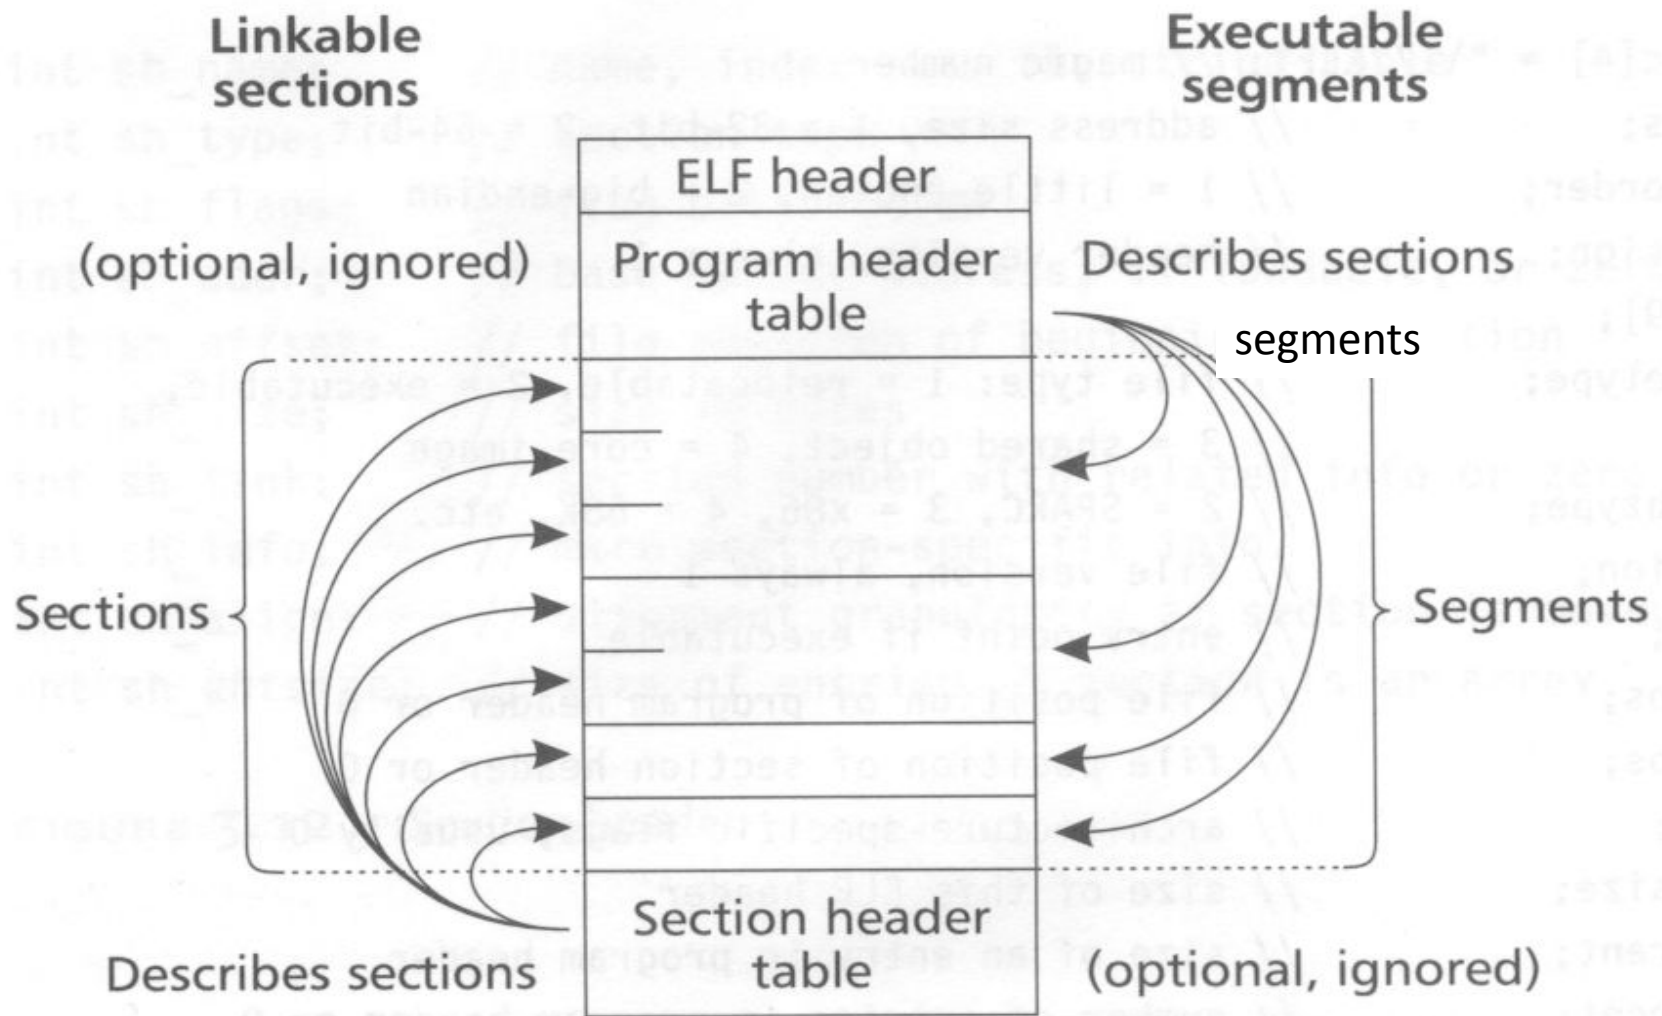


FIGURE 3.10 • Two views of an ELF file.

ELF File Format

- An ELF file has:
 - A header
 - A program header
 - A section header
 - Sections
 - More stuff which matters less (string table, debug symbols, ...)

Sections are used by the linker (compile time)

Segments are used by the loader (runtime)

ELF Structure

- A single segment usually consist of several sections. E.g., a loadable read-only segment could contain sections for executable code, read-only data, and symbols for the dynamic linker.
- Relocatable files have section header tables. Executable files have program header tables. Shared object files have both.
- Sections are intended for further processing by a linker, while the segments are intended to be mapped into memory.

An Example C Program

```
#include <stdio.h>  
int main() {  
    printf("hello world!\n");  
}
```

Save the code in a file named test.c, and then compile the source code into a *.o file and an executable with the commands below.

```
$ gcc test.c -c
```

```
$ gcc test.o -o test
```

We can check if a file is ELF file by the file command as shown below.

```
$ file test.o
```


Program segments/sections

```
$ size /usr/bin/test.o
```

text	data	bss	dec	hex
245138	3696	2832	251666	3d712

```
filename  
/usr/bin/cc
```

```
$ $ readelf -h test
```

The -h option means to display the ELF file header. The Magic number is used to indicate the file is ELF file.

ELF Header

- The Elf header is always at offset zero of the file.
- The program header table and the section header table's offset in the file are defined in the ELF header.
- The header is decodable even on machines with a different byte order from the file's target architecture.

\$ *readelf -h test.o*

The -h option means to display the ELF file header. The Magic number is used to indicate the file is ELF file.

```

char magic[4] = "\177ELF"; // magic number
char class;                // address size, 1 = 32-bit, 2 = 64-bit
char byteorder;            // 1 = little-endian, 2 = big-endian
char hversion;            // header version, always 1
char pad[9];
short filetype;            // file type: 1 = relocatable, 2 = executable,
                           // 3 = shared object, 4 = core image
short archtype;            // 2 = SPARC, 3 = x86, 4 = 68K, etc.
int fversion;              // file version, always 1
int entry;                 // entry point if executable
int phdrpos;               // file position of program header or 0
int shdrpos;               // file position of section header or 0
int flags;                 // architecture-specific flags, usually 0
short hdrsize;             // size of this ELF header
short phdrent;             // size of an entry in program header
short phdrcnt;             // number of entries in program header or 0
short shdrent;             // size of an entry in section header
short shdrcnt;             // number of entries in section header or 0
short strsec;              // section number that contains section name strings

```

FIGURE 3.11 • ELF header.

Display ELF header: readelf -h

ELF Header:

```
Magic:      7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
Class:                                ELF64
Data:                                2's complement, little
    endian
Version:                                1 (current)
OS/ABI:                                UNIX - System V
ABI Version:                            0
Type:                                EXEC (Executable file)
Machine:                                Advanced Micro Devices
    X86-64
Version:                                0x1
Entry point address:                    0x400440
Start of program headers:                64 (bytes into file)
Start of section headers:                4424 (bytes into file)
Flags:                                0x0
Size of this header:                    64 (bytes)
Size of program headers:                56 (bytes)
Number of program headers:                9
Size of section headers:                64 (bytes)
Number of section headers:                31
Section header string table index:      28
```

ELF File Format

Linking View




Execution View





ELF Files: Dual Views

Linking View




Execution View





Linker Overview

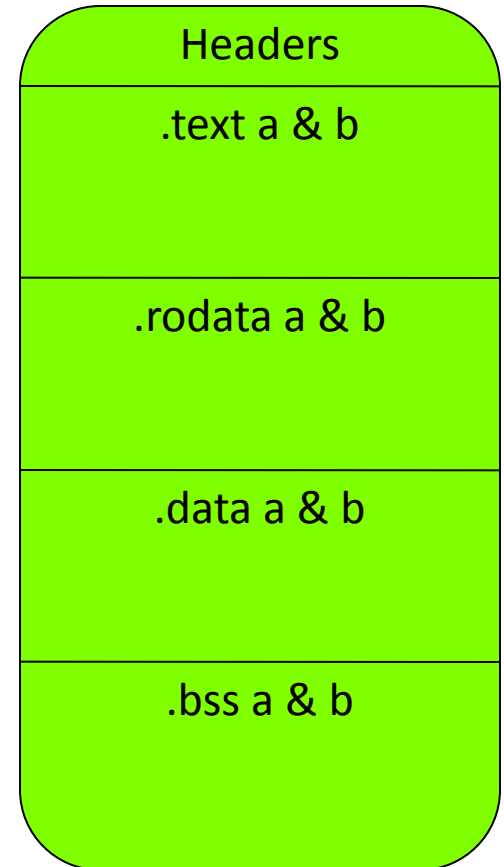
ELF object file a



ELF object file b

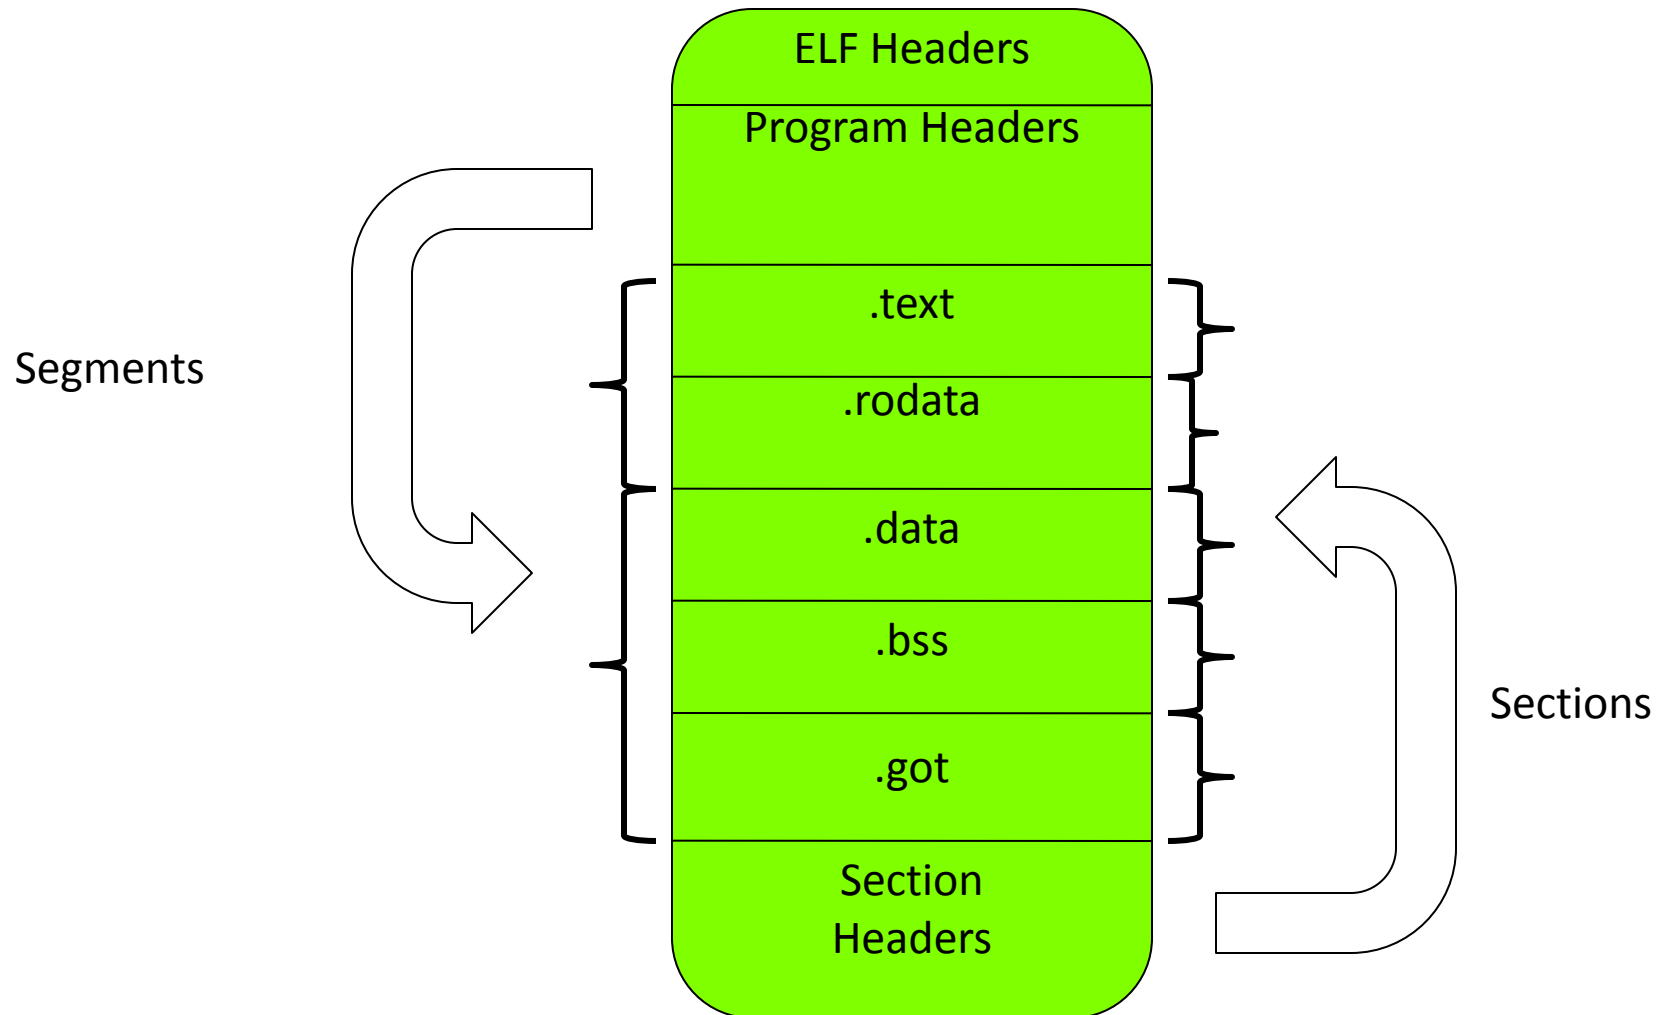


Executable or library



Headers Overview

Executable or
library



Relocatable Files

- A relocatable or shared object file is a collection of sections.
- Each section contains a single type of information, such as program code, read-only data, or read/write data, relocation entries, or symbols.
- Every symbol's address is defined relative to a section.
 - Therefore, a procedure's entry point is relative to the program code section that contains that procedure's code.

Section Header

```
int sh_name;      // name, index into the string table
int sh_type;      // section type
int sh_flags;     // flag bits, below
int sh_addr;      // base memory address, if loadable, or zero
int sh_offset;    // file position of beginning of section
int sh_size;      // size in bytes
int sh_link;      // section number with related info or zero
int sh_info;      // more section-specific info
int sh_align;     // alignment granularity if section is moved
int sh_entsize;   // size of entries if section is an array
```

FIGURE 3.12 • Section header.

Types in Section Header

- PROGBITS: This holds program contents including code, data, and debugger information.
- NOBITS: Like PROGBITS. However, it occupies no space.
- SYMTAB : These hold symbol table.
- STRTAB: This is a string table, like the one used in a.out.
- REL and RELA: These hold relocation information.
- DYNAMIC and HASH: This holds information related to dynamic linking.

Flags in Section Header

- WRITE: This section contains data that is writable during process execution.
- ALLOC: This section occupies memory during process execution.
- EXECINSTR: This section contains executable machine instructions.

Various Sections

- `.text`:
 - This section holds executable instructions of a program.
 - Type: PROGBITS
 - Flags: ALLOC + EXECINSTR
 - `$ readelf -S vipi | grep '.text'`
- `.data`:
 - This section holds initialized data that contributes to the program's image.
 - Type: PROGBITS
 - Flags: ALLOC + WRITE
- `.rodata`:
 - This section holds read-only data.
 - Type: PROGBITS
 - Flags: ALLOC
 - `$objdump -s test.o`

Various Sections

- `.bss` :
 - This section holds uninitialized data that contributed to the program's image. By definition, the system will initialize the data with zero when the program begins to run.
 - Type: NOBITS
 - Flags: ALLOC + WRITE
- `.rel.text`, `.rel.data`, and `.rel.rodata`:
 - These contain the relocation information for the corresponding text or data sections.
 - Type: REL
 - Flags: ALLOC is turned on if the file has a loadable segment that includes relocation.
- `.symtab`:
 - This section hold a symbol table.
- `.strtab`:
 - This section holds strings.

Various Sections

- `.init:`
 - This section holds executable instructions that contribute to the process initialization code.
 - Type: `PROGBITS`
 - Flags: `ALLOC + EXECINSTR`
- `.fini:`
 - This section hold executable instructions that contribute to the process termination code.
 - Type: `PROGBITS`
 - Flags: `ALLOC + EXECINSTR`
- C does not need these two sections. However, C++ needs them.

Various Sections

- `.interp`:
 - This section holds the pathname of a program interpreter.
 - Type: `ALLOC`
 - Flags: `PROGBITS`
 - If this section is present, rather than running the program directly, the system runs the interpreter and passes it the elf file as an argument.
 - For many years (used in `a.out`), UNIX has had self-running interpreted text files, using
 - `#!/bin/csh` as the first line of the file.
 - Elf extends this facility to interpreters that run nontext programs.
 - In practice, this is used to run the run-time dynamic linker to load the program and to link in any required shared libraries.

Various Sections

- `.debug:`
 - This section holds symbolic debugging information.
 - Type: PROGBIT
- `.line:`
 - This section holds line number information for symbolic debugging, which describes the correspondence between the program source and the machine code (ever used gdb?)
 - Type: PROGBIT
- `.comment`
 - This section may store extra information.
- `.got:`
 - This section holds the global offset table.
 - Related to shared library.
 - Type: PROGBIT
- `.plt:`
 - This section holds the procedure linkage table.
 - Type: PROGBIT
- `.note:`
 - This section contains some extra information.

Various Sections

- `.debug:`
 - This section holds symbolic debugging information.
 - Type: PROGBIT
- `.line:`
 - This section holds line number information for symbolic debugging, which describes the correspondence between the program source and the machine code (ever used gdb?)
 - Type: PROGBIT
- `.comment`
 - This section may store extra information.
- `.got:`
 - This section holds the global offset table.
 - Related to shared library.
 - Type: PROGBIT
- `.plt:`
 - This section holds the procedure linkage table.
 - Type: PROGBIT
- `.note:`
 - This section contains some extra information.

ELF header	} (not considered sections)
(segment table)	
.text	
.data	
.rodata	
.bss	
.sym	
.rel.text	
.rel.data	
.rel.rodata	
.line	
.debug	
.strtab	
Section table	(not considered a section)

A typical relocatable file.

FIGURE 3.14 • Sample relocatable ELF file.

String Table

- Like the format used in a.out.
- String table sections hold null-terminated character sequences, commonly called strings.
- The object file uses these strings to represent symbol and section names.
- We use an index into the string table section to reference a string.
- The reason why we separate symbol names from symbol tables is that in C or C++, there is no limitation on the length of a symbol.

String Table

Index	+0	+1	+2	+3	+4	+5	+6	+7	+8	+9
0	\0	n	a	m	e	.	\0	V	a	r
10	i	a	b	l	e	\0	a	b	l	e
20	\0	\0	x	x	\0					

\$ readelf -p '.symtab' test.o

Index	String
0	<i>none</i>
1	name.
7	Variable
11	able
16	able
24	<i>null string</i>

Symbol Table

- An object file's symbol table holds information needed to locate and relocate a program's symbolic definition and references.
- A symbol table index is a subscript into this array.

```
$ readelf -s test.o
```

Symbol Table

```
int name;      // position of name string in string table
int value;     // symbol value, section relative in reloc,
               // absolute in executable
int size;      // object or function size
char type:4;   // data object, function, section, or special-case file
char bind:4;   // local, global, or weak
char other;    // spare
short sect;    // section number, ABS, COMMON, or UNDEF
```

e.g., int, double

If a definition is available for an undefined weak symbol, the linker will use it. Otherwise, the value defaults to 0.

FIGURE 3.13 • ELF symbol table.

The section relative to which the symbol is defined. (e.g., the function entry points are defined relative to .text)

Relocation Table

- Relocation is the process of connecting symbolic references with symbolic definitions.
- Relocatable files must have information that describes how to modify their section contents.
- A relocation table consists on many relocation structures.

- *\$ readelf -l test.o* to read program Headers
- *\$ readelf -S test.o* view the section headers and program headers with readelf command.

Executable Files

- An executable file usually has only a few segments.
E.g.,
 - A read-only one for the code.
 - A read-only one for read-only data.
 - A read/write one for read/write data.
- All of the loadable sections are packed into the appropriate segments so that the system can map the file with just one or two operations.
 - E.g., If there is a .init and .fini sections, those sections will be put into the read-only text segment.
 - Segment are Page Aligned

Segments

- Executable and shared object files contain segments. A segment is a grouping of one or more sections in the object file. For example, the *code segment* includes the .text and .rodata sections while the *data segment* contains the .bss and .data sections.
- It Defines:
- Permission applied in segments
- Size & Section Information
- Sections with identical permission can mapped to one segments
- **\$ readelf -l test.o**

Dynamic linking

- Dynamic linking allows to resolve library functions at runtime
- Relocatable File: Holds code and data suitable for linking with other object files to create an executable or shared object file
 - Executable File: Holds a program suitable for execution
 - Shared Object File: Holds code and data suitable for linking in two contexts:
 - The **Link Editor** may process it with other relocatable and shared object files to create another object file
 - The **Dynamic Linker** combines it with an executable file and other shared objects to create a process image

Program Header

```
int type;      // loadable code or data, dynamic linking info, etc.
int offset;    // file offset of segment
int virtaddr;  // virtual address to map segment
int physaddr;  // physical address, not used
int filesize;  // size of segment in file
int memsize;   // size of segment in memory (bigger if contains bss)
int flags;     // Read, Write, Execute bits
int align;     // required alignment, invariably hardware page size
```

FIGURE 3.15 • ELF program header.

The Types in Program Header

- This field tells what kind of segment this array element describes:
 - PT_LOAD: This segment is a loadable segment.
 - PT_DYNAMIC: This array element specifies dynamic linking information.
 - PT_INTERP: This element specified the location and size of a null-terminated path name to invoke as an interpreter.

Executable File Example

	<i>File offset</i>	<i>Load address</i>	<i>Type header</i>
ELF header	0	0x80000000	
Program header	0x40	0x80000040	
Read-only text (size 0x4500)	0x100	0x8000100	LOAD, read/execute
Read/write data (file size 0x2200, memory size 0x3500)	0x4600	0x8005600	LOAD, read/write/ execute

Nonloadable information and optional section headers

FIGURE 3.16 • ELF loadable segments.

Elf Linking

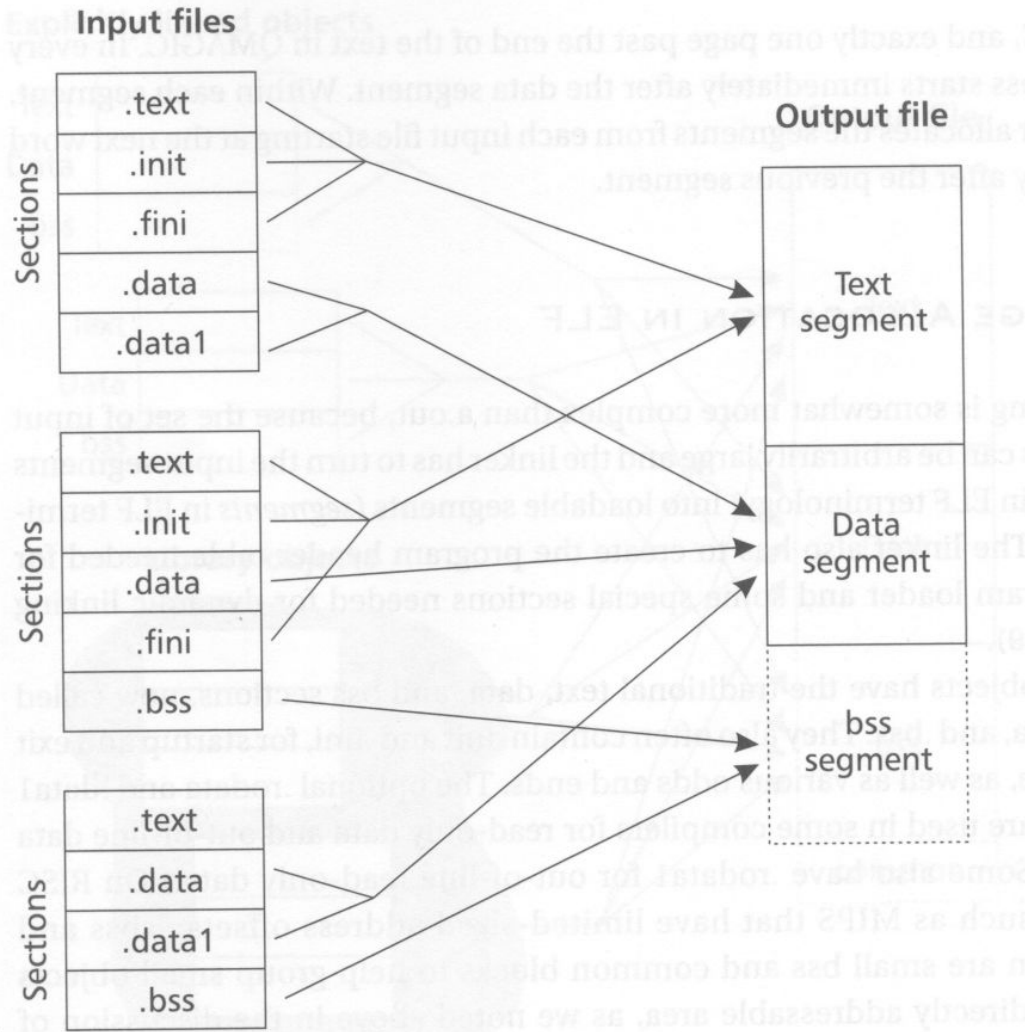


FIGURE 4.9 • ELF linking.

Elf File Trace

(We can use the objdump or nm command)

Program Header

Program Header:

PHDR	off	0x00000034	vaddr	0x08048034	paddr	0x08048034	align	2**2
	filesz	0x000000c0	memsz	0x000000c0	flags	r-x		
INTERP	off	0x000000f4	vaddr	0x080480f4	paddr	0x080480f4	align	2**0
	filesz	0x00000019	memsz	0x00000019	flags	r--		
LOAD	off	0x00000000	vaddr	0x08048000	paddr	0x08048000	align	2**12
	filesz	0x000000564	memsz	0x000000564	flags	r-x		
LOAD	off	0x000000564	vaddr	0x08049564	paddr	0x08049564	align	2**12
	filesz	0x0000000a8	memsz	0x0000000cc	flags	rw-		
DYNAMIC	off	0x00000059c	vaddr	0x0804959c	paddr	0x0804959c	align	2**2
	filesz	0x000000070	memsz	0x000000070	flags	rw-		
NOTE	off	0x000000110	vaddr	0x08048110	paddr	0x08048110	align	2**2
	filesz	0x000000018	memsz	0x000000018	flags	r--		

Dynamic Section

Dynamic Section:

NEEDED	libc.so.4
INIT	0x8048390
FINI	0x8048550
HASH	0x8048128
STRTAB	0x80482c8
SYMTAB	0x80481b8
STRSZ	0xad
SYMENT	0x10
DEBUG	0x0
PLTGOT	0x8049584
PLTRELSZ	0x18
PLTREL	0x11
JMPREL	0x8048378

Need to link this shared library
for printf()

Section Header

Sections:

Idx	Name	Size	VMA	LMA	File off	Algn
0	.interp	00000019	080480f4	080480f4	000000f4	2**0
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
1	.note.ABI-tag	00000018	08048110	08048110	00000110	2**2
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
2	.hash	00000090	08048128	08048128	00000128	2**2
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
3	.dynsym	00000110	080481b8	080481b8	000001b8	2**2
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
4	.dynstr	000000ad	080482c8	080482c8	000002c8	2**0
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
5	.rel.plt	00000018	08048378	08048378	00000378	2**2
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
6	.init	0000000b	08048390	08048390	00000390	2**2
	CONTENTS, ALLOC, LOAD, READONLY, CODE					
7	.plt	00000040	0804839c	0804839c	0000039c	2**2
	CONTENTS, ALLOC, LOAD, READONLY, CODE					
8	.text	00000174	080483dc	080483dc	000003dc	2**2
	CONTENTS, ALLOC, LOAD, READONLY, CODE					

Section Header (cont'd)

9	.fini	00000006	08048550	08048550	00000550	2**2
		CONTENTS, ALLOC, LOAD, READONLY, CODE				
10	.rodata	0000000e	08048556	08048556	00000556	2**0
		CONTENTS, ALLOC, LOAD, READONLY, DATA				
11	.data	0000000c	08049564	08049564	00000564	2**2
		CONTENTS, ALLOC, LOAD, DATA				
12	.eh_frame	00000004	08049570	08049570	00000570	2**2
		CONTENTS, ALLOC, LOAD, DATA				
13	.ctors	00000008	08049574	08049574	00000574	2**2
		CONTENTS, ALLOC, LOAD, DATA				
14	.dtors	00000008	0804957c	0804957c	0000057c	2**2
		CONTENTS, ALLOC, LOAD, DATA				
15	.got	00000018	08049584	08049584	00000584	2**2
		CONTENTS, ALLOC, LOAD, DATA				
16	.dynamic	00000070	0804959c	0804959c	0000059c	2**2
		CONTENTS, ALLOC, LOAD, DATA				
17	.bss	00000024	0804960c	0804960c	0000060c	2**2
		ALLOC				
18	.stab	000001bc	00000000	00000000	0000060c	2**2
		CONTENTS, READONLY, DEBUGGING				
19	.stabstr	00000388	00000000	00000000	000007c8	2**0
		CONTENTS, READONLY, DEBUGGING				
20	.comment	000000c8	00000000	00000000	00000b50	2**0

Symbol Table

SYMBOL TABLE:

080480f4	1	d	.interp	00000000
08048110	1	d	.note.ABI-tag	00000000
08048128	1	d	.hash	00000000
080481b8	1	d	.dynsym	00000000
080482c8	1	d	.dynstr	00000000
08048378	1	d	.rel.plt	00000000
08048390	1	d	.init	00000000
0804839c	1	d	.plt	00000000
080483dc	1	d	.text	00000000
08048550	1	d	.fini	00000000
08048556	1	d	.rodata	00000000
08049564	1	d	.data	00000000
08049570	1	d	.eh_frame	00000000
08049574	1	d	.ctors	00000000
0804957c	1	d	.dtors	00000000
08049584	1	d	.got	00000000
0804959c	1	d	.dynamic	00000000

Symbol Table (cont'd)

- 0804960c 1 d .bss 00000000
- 00000000 1 d .stab 00000000
- 00000000 1 d .stabstr 00000000
- 00000000 1 d .comment 00000000
- 00000000 1 d .note 00000000
- 00000000 1 d *ABS* 00000000
- 00000000 1 d *ABS* 00000000
- 00000000 1 d *ABS* 00000000
- 00000000 1 df *ABS* 00000000 crtstuff.c
- 08048460 1 .text 00000000 gcc2_compiled.
- 08049568 1 O .data 00000000 p.3
- 0804957c 1 O .dtors 00000000 __DTOR_LIST__
- 0804956c 1 O .data 00000000 completed.4
- 08048460 1 F .text 00000000 __do_global_dtors_aux
- 08049570 1 O .eh_frame 00000000 __EH_FRAME_BEGIN__

Symbol Table (cont'd)

080484b4	l	F .text	00000000	fini_dummy
0804960c	l	O .bss	00000018	object.11
080484bc	l	F .text	00000000	frame_dummy
080484e0	l	F .text	00000000	init_dummy
08049570	l	O .data	00000000	force_to_data
08049574	l	O .ctors	00000000	__CTOR_LIST__
00000000	l	df *ABS*	00000000	crtstuff.c
08048520	l	.text	00000000	gcc2_compiled.
08048520	l	F .text	00000000	__do_global_ctors_aux
08049578	l	O .ctors	00000000	__CTOR_END__
08048548	l	F .text	00000000	init_dummy
08049570	l	O .data	00000000	force_to_data
08049580	l	O .dtors	00000000	__DTOR_END__
08049570	l	O .eh_frame	00000000	__FRAME_END__
00000000	l	df *ABS*	00000000	p10.c
080483ac		F *UND*	00000031	printf
0804959c	g	O *ABS*	00000000	_DYNAMIC
08048550	g	O *ABS*	00000000	_etext
08048390	g	F .init	00000000	_init
08049624	g	O .bss	00000004	environ
00000000	w	*UND*	00000000	__deregister_frame_info
08049630	g	O *ABS*	00000000	end
08049628	g	O .bss	00000004	xx

Symbol Table (cont'd)

08049564	g	O	.data	00000004	__progrname
080483dc	g	F	.text	00000083	_start
0804960c	g	O	*ABS*	00000000	__bss_start
080484e8	g	F	.text	00000038	main
08048550	g	F	.fini	00000000	_fini
0804962c	g	O	.bss	00000004	yy
080483bc		F	*UND*	00000070	atexit
0804960c	g	O	*ABS*	00000000	_edata
08049584	g	O	*ABS*	00000000	_GLOBAL_OFFSET_TABLE_
08049630	g	O	*ABS*	00000000	_end
080483cc		F	*UND*	0000005b	exit
00000000	w		*UND*	00000000	__register_frame_info

Dynamic Symbol Table

DYNAMIC SYMBOL TABLE:

080483ac		DF	*UND*	00000031	printf
0804959c	g	DO	*ABS*	00000000	__DYNAMIC
08048550	g	DO	*ABS*	00000000	__etext
08048390	g	DF	.init	00000000	__init
08049624	g	DO	.bss	00000004	environ
00000000	w	D	*UND*	00000000	__deregister_frame_info
08049630	g	DO	*ABS*	00000000	end
08049564	g	DO	.data	00000004	__progrname
0804960c	g	DO	*ABS*	00000000	__bss_start
08048550	g	DF	.fini	00000000	__fini
080483bc		DF	*UND*	00000070	atexit
0804960c	g	DO	*ABS*	00000000	__edata
08049584	g	DO	*ABS*	00000000	__GLOBAL_OFFSET_TABLE__
08049630	g	DO	*ABS*	00000000	__end
080483cc		DF	*UND*	0000005b	exit
00000000	w	D	*UND*	00000000	__register_frame_info

Dynamic Relocation Table

- DYNAMIC RELOCATION RECORDS

• OFFSET	TYPE	VALUE
• 08049590	R_386_JUMP_SLOT	printf
• 08049594	R_386_JUMP_SLOT	atexit
• 08049598	R_386_JUMP_SLOT	exit

Library

Library is a collection of resources used to develop software. These may include pre-written code and subroutines, classes, values or type specifications.

Libraries contain code and data that provide services to independent programs.

Adv- sharing and changing of code and data in a modular fashion, and eases the distribution of the code and data.

Some executables are both standalone programs and libraries, but most libraries are not executable. Executables and libraries make references known as links to each other through the process known as linking, which is typically done by a linker.

Shared Libraries

shared library or shared object is a file that is intended to be shared by executable files and further shared objects files.

Modules used by a program are loaded from individual shared objects into memory at load time or run time, rather than being copied by a linker when it creates a single monolithic executable file for the program.

Shared Libraries

Shared libraries remove the common library routines from the executable file

Maintaining a single copy of the library routine somewhere in memory that all processes reference.

This reduces the size of each executable file (But may add some runtime overhead, either when the program is first executed or the first time each shared library function is called)

Another advantage of shared libraries is that library functions can be replaced with new versions without having to relink edit every program that uses the library.

```
$ cc -static hello1.c          prevent gcc from using shared libraries
```

```
$ ls -l a.out  
-rwxrwxr-x 1 sar      475570 Feb 18 23:17 a.out
```

```
$ size a.out  
   text    data    bss    dec    hex    filename  
375657   3780   3220 382657 5d6c1   a.out
```

If we compile this program to use shared libraries, the text and data sizes of the executable file are greatly decreased:

```
$ cc hello1.c                  gcc defaults to use shared libraries
```

```
$ ls -l a.out  
-rwxrwxr-x 1 sar      11410 Feb 18 23:19 a.out
```

```
$ size a.out  
   text    data    bss    dec    hex    filename  
   872     256     4  1132    46c    a.out
```

<http://tldp.org/HOWTO/Program-Library-HOWTO/shared-libraries.html>